

SMC: Exercises set IV

Niharika

Sep 22, 2018

This document provides solution for the Exercises set IV given at http://www.it.uu.se/research/systems_and_control/education/2017/smc/homework/SMC2017_exercises4.pdf.

IV.1 Particle Metropolis-Hastings

Consider the following state-space model:

$$\begin{aligned}x_t &= \cos(\theta * x_{t-1}) + v_t, & v_t &\sim N(0, 1) \\y_t &= x_t + e_t, & e_t &\sim N(0, 1) \\x_0 &\sim N(0, 1)\end{aligned}$$

Generate y assuming $\theta = 1$.

```
T <- 50
N <- 100

# simulate data
simulateData <- function(param, T)
{
  theta <- param[1]
  sigma_v <- param[2]
  sigma_e <- param[3]

  x <- rep(0, T)
  y <- rep(0, T)

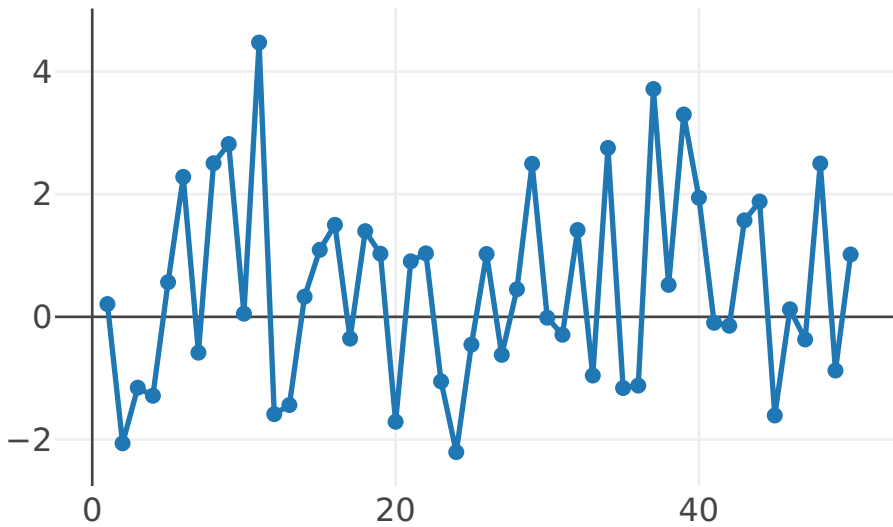
  x0 <- 0 # state 0
  x[1] <- cos(theta*x0) + rnorm(1, 0, sigma_v)
  y[1] <- x[1] + rnorm(1, 0, sigma_e)

  for(t in 2:T)
  {
    x[t] <- cos(theta*x[t-1]) + rnorm(1, 0, sigma_v)
    y[t] <- x[t] + rnorm(1, 0, sigma_e)
  }

  return(list(x=x, y=y))
}

param <- c(1, 1, 1)
simulatedData <- simulateData(param, T)

plot_ly(x = c(1:T), y = simulatedData$y, name = 'simulated states', type = 'scatter', mode = 'lines+markers')
```



```
bootstrapParticleFilter <- function(param, y, N) {
  T <- length(y)
  theta <- param[1]
  sigma_v <- param[2]
  sigma_e <- param[3]

  # Initialize variables
  particles <- matrix(0, nrow = N, ncol = T)
  weights <- matrix(1, nrow = N, ncol = T)
  normalisedWeights <- matrix(0, nrow = N, ncol = T)
  xHat <- rep(0, T)
  logLikelihood <- 0

  #state at t=0
  x0 <- rnorm(N, 0, 1)

  particles[,1] <- x0
  normalisedWeights[, 1] = 1 / N

  for (t in 2:T) {
    # Resample
    newAncestors <- sample(N, replace = TRUE, prob = normalisedWeights[, t - 1])

    # Propagate
    particles[, t] <- cos(theta * particles[newAncestors, t-1]) +
      rnorm(N, 0, 1)

    #Likelihood
    weights[, t] <- dnorm(y[t], mean = particles[, t], sd = 1, log = TRUE)

    max_weight <- max(weights[, t])
    weights[, t] <- exp(weights[, t] - max_weight)
    normalisedWeights[, t] <- weights[, t] / sum(weights[, t])

    # Estimate the state
    xHat[t] <- mean(particles[, t])

    # accumulate the log-likelihood
    logLikelihood = logLikelihood + max_weight +
      log(sum(weights)) - log(N)
  }
}
```

```

list(xHat = xHat,
     logLikelihood = logLikelihood,
     particles = particles,
     weights = normalisedWeights)
}

particleMH <- function(y, N, iter)
{
  theta <- rep(0, iter)
  theta_proposed <- rep(0, iter)
  logLikelihood <- rep(0, iter)
  logLikelihood_proposed <- rep(0, iter)

  # initialize
  theta[1] <- rnorm(1, 0, 1)
  param <- c(theta[1], 1, 1)
  result <- bootstrapParticleFilter(y, param, N)
  logLikelihood[1] <- result$logLikelihood

  for (k in 2:iter) {
    # Propose a new parameter
    theta_proposed[k] <- theta[k - 1] + rnorm(1)

    param <- c(theta_proposed[k], 1, 1)
    result <- bootstrapParticleFilter(y, param, N)
    logLikelihood_proposed[k] <- result$logLikelihood

    # Compute the acceptance probability
    numerator = logLikelihood_proposed[k] + dnorm(theta_proposed[k], log = TRUE)
    denominator = logLikelihood[k-1] + dnorm(theta[k-1], log = TRUE)

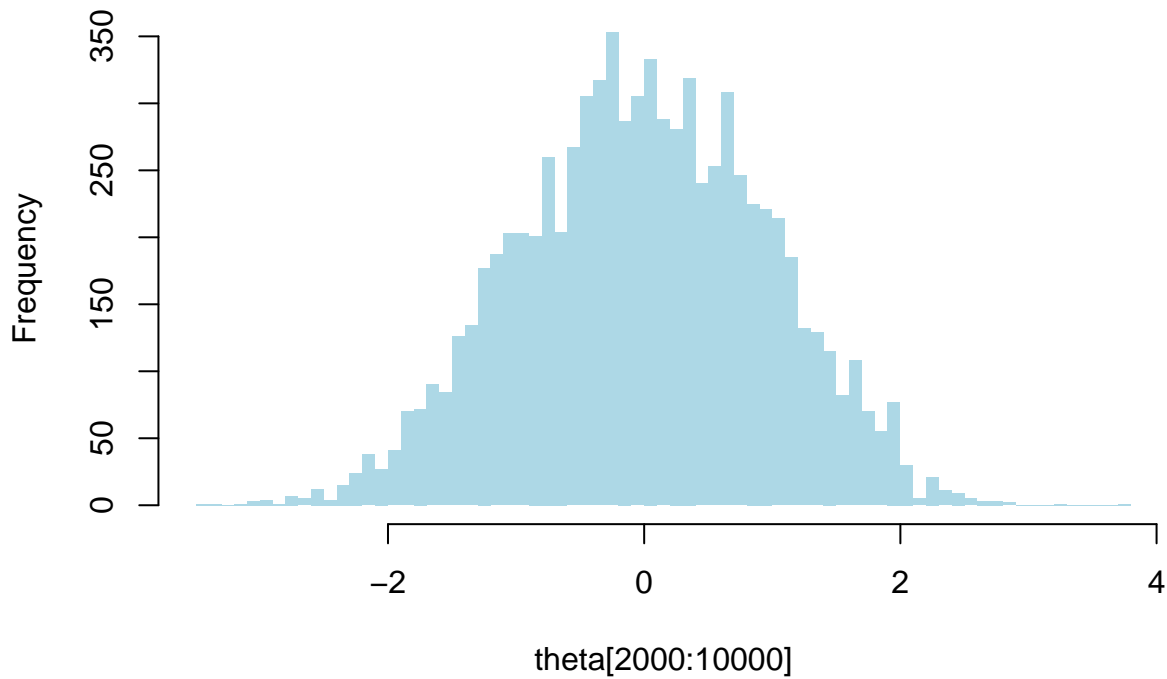
    # Compute acceptance probability
    acc_prob = exp(numerator - denominator)
    acc_prob = min(1, acc_prob)

    u <- runif(1)
    if (u < acc_prob) {
      # Accept the parameter
      theta[k] <- theta_proposed[k]
      logLikelihood[k] <- logLikelihood_proposed[k]
    } else {
      # Reject the parameter
      theta[k] <- theta[k - 1]
      logLikelihood[k] <- logLikelihood[k - 1]
    }
  }
  return(theta)
}

theta <- particleMH(y = simulatedData$y, N, 10000)
hist(theta[2000:10000], breaks = 100, lty=0,
     main = "Posterior distribution of theta", col = "lightblue")

```

Posterior distribution of theta



IV.2 Conditional particle

Fully adapted particle filter is transformed into conditional particle filter.

```
T <- 100
N <- 100
# simulate data
simulateData <- function(param, T)
{
  sigma_v <- param[1]
  sigma_e <- param[2]

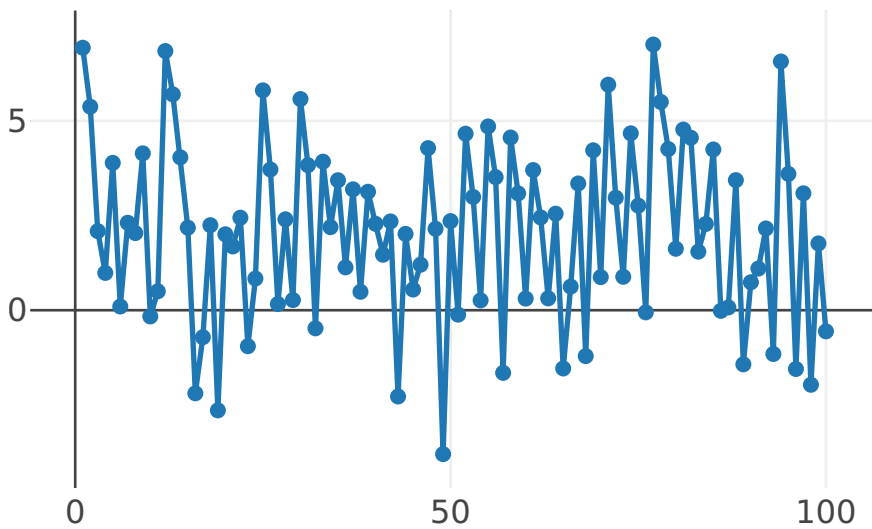
  x <- rep(0, T)
  y <- rep(0, T)

  #initial state
  x0 <- 0
  x[1] <- 2*cos(x0)^2 + rnorm(1, 0, sigma_v)
  y[1] <- 2 * x[1] + rnorm(1, 0, sigma_e)

  for(t in 2:T)
  {
    x[t] <- 2*cos(x[t-1])^2 + rnorm(1, 0, sigma_v)
    y[t] <- 2*x[t] + rnorm(1, 0, sigma_e)
  }

  return(list(x=x, y=y))
}

param <- c(1, .1)
simulatedData <- simulateData(param, T)
plot_ly(x = c(1:T), y = simulatedData$y, name = 'Observed data', type = 'scatter', mode = 'lines+markers')
```



Conditions Fully Adapted Particle Filter#####

```
conditionalAPF <- function(x_ref, y, N)
{
  T <- length(y)

  # Initialize variables
  particles <- matrix(0, nrow = N, ncol = T)
  weights <- matrix(1, nrow = N, ncol = T)
  normalisedWeights <- matrix(0, nrow = N, ncol = T)
  xHat <- rep(0, T)
  logLikelihood <- 0

  #state at t=0
  x0 <- rnorm(N, 0, 1)
  # Replace last state with state from reference trajectory
  x0[1] = x_ref[1]

  particles[,1] <- x0
  normalisedWeights[, 1] = 1 / N

  for (t in 2:T) {
    # compute resampling weights
    weights[, t] <- dnorm(y[t], mean = 2*cos(particles[, t-1])^2,
                          sd = sqrt(4.01), log = TRUE)
    max_weight <- max(weights[, t])
    weights[, t] <- exp(weights[, t] - max_weight)
    normalisedWeights[, t] <- weights[, t] / sum(weights[, t])

    newAncestors <- sample(N, replace = TRUE, prob = normalisedWeights[, t])

    # Propagate
    particles[, t] <- 0.1/sqrt(4.01)*rnorm(N, 0, 1) +
      (2/4.01) * y[t] + .01/4.01 *
      cos(particles[newAncestors, t-1])^2

    # Replace last sample with reference trajectory
    newAncestors[N] = N
    particles[N, t] = x_ref[t]
  }

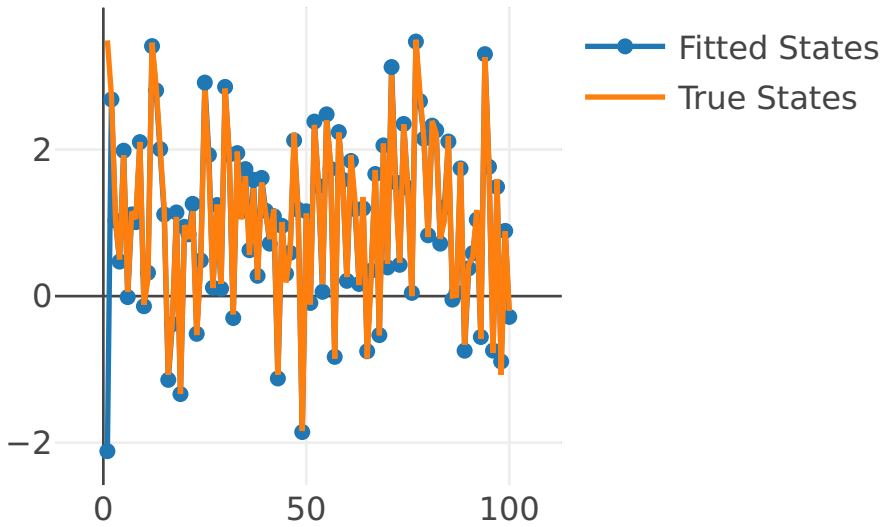
  b = sample(1:N, 1)
  return(particles[b, ])
}
```

```

}

x = conditionalAPF(x_ref = simulatedData$x, y = simulatedData$y, N)
p <- plot_ly(x = c(1:T), y = x, name = 'Fitted States', type = 'scatter', mode = 'lines+markers')
add_lines(p, x = c(1:T), y = simulatedData$x, name = 'True States', type = 'scatter', mode = 'lines+markers')

```



IV.3 Conditional importance sampling

- (a) With $N=2$ particles, sample from $\pi(x) = \mathcal{N}(x|1, 1)$ by using conditional importance sampling with the proposal $q(x) = \mathcal{N}(x|0, 1)$

```

condImpSamplingKernel <- function(x_ref, N=2)
{
  # Sample from proposal
  particles <- rnorm(N)

  # Set the last sample to the reference
  particles[N] = x_ref

  # Calculate weights
  w = dnorm(particles, mean = 1, sd = 1, log = TRUE) -
      dnorm(particles, mean = 0, sd = 1, log = TRUE)
  w = w - max(w)
  w = exp(w) / sum(exp(w))

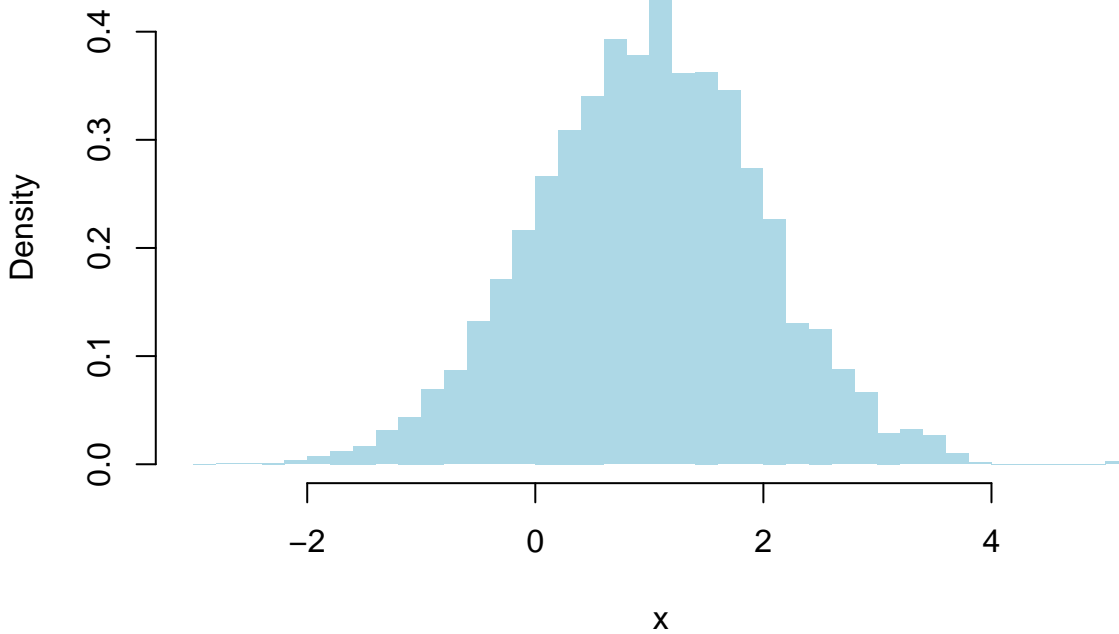
  return(particles[sample(N, size = 1, prob = w)])
}

condImpSampling <- function(iter = 100, N = 2)
{
  # Initialize variables
  particles <- rep(0, iter)
  particles[1] <- rnorm(1)
  for (i in 2:iter) {
    particles[i] = condImpSamplingKernel(particles[i-1], N = N)
  }
  return(particles)
}

x <- condImpSampling(50000)
hist(x, breaks = 50, prob = TRUE, col = "lightblue", lty=0)

```

Histogram of x



IV.4 An SMC sampler for localization

(a) Simulate M independent measurements from the model

$$\begin{aligned} y_m^1 &= x_0^1 + n_m^1 b_m^1 \\ y_m^2 &= x_0^2 + n_m^2 b_m^2 \end{aligned}$$

where

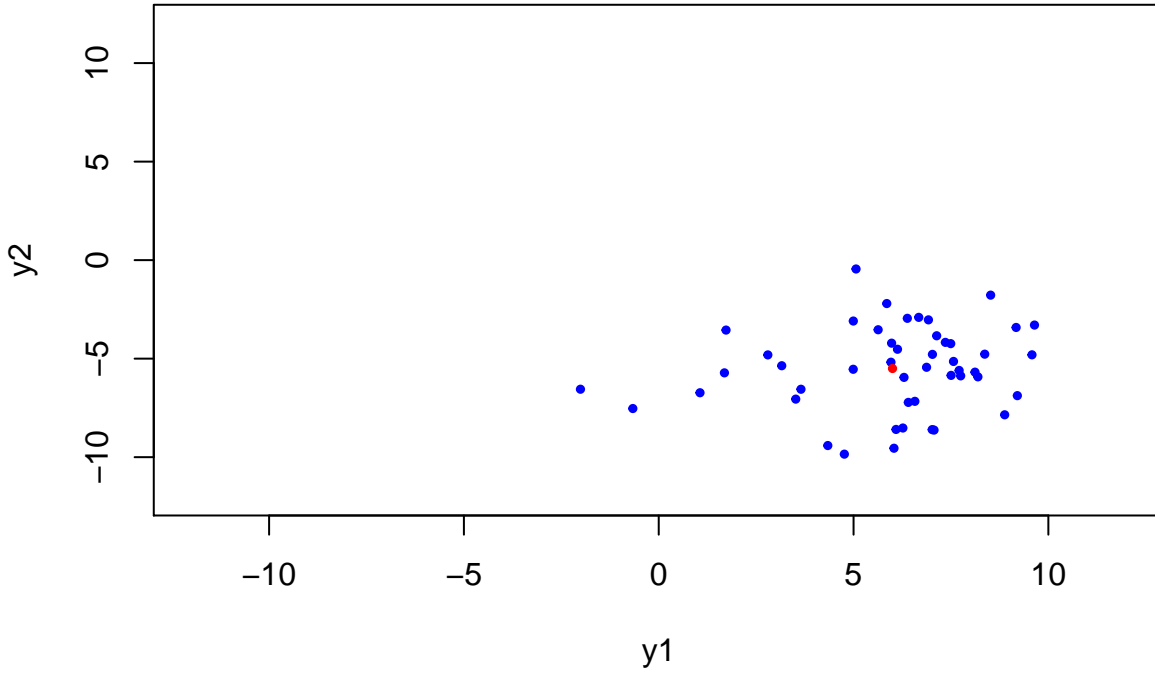
$$\begin{aligned} m &= 1, 2, \dots, M \\ x_0 &= (x_0^1, x_0^2) \\ n_m^1, n_m^2 &\sim \text{Exp}(2) \\ \mathbb{P}(b_m^1 = 1) &= \mathbb{P}(b_m^1 = -1) = \frac{1}{2} \\ \mathbb{P}(b_m^2 = 1) &= \mathbb{P}(b_m^2 = -1) = \frac{1}{2} \end{aligned}$$

```
M = 50
x0 = c(6, -5.5)
n1 = rexp(M, 1/2)
n2 = rexp(M, 1/2)
n <- cbind(n1, n2)

b1 = sample(c(-1,1), size = M, replace = TRUE, prob = c(0.5, 0.5))
b2 = sample(c(-1,1), size = M, replace = TRUE, prob = c(0.5, 0.5))
b <- cbind(b1, b2)

y1 = x0[1] + n1*b1
y2 = x0[2] + n2*b2

y = cbind(y1, y2)
plot(y, xlim = c(-12,12), ylim = c(-12,12), cex = .5, col = "blue", pch = 19)
lines(x0[1], x0[2], col = "red", type = "p", cex = .5, pch = 19)
```



(b) Likelihood The components of y_m are independent

$$p(y_m | x_0) = p(y_m^1 | x_0^1) \cdot p(y_m^2 | x_0^2)$$

where

$$p(y_m^i | x_0^i) = \frac{1}{4} \exp\left(-\frac{\text{abs}(y_m^i - x_0^i)}{2}\right)$$

```
logLikeliHood <- function(x, y)
{
  sum_ll <- c(0, 0)

  for (m in 1:M) {
    sum_ll <- sum_ll + log(0.25) + 0.5 * c(-1, -1) * abs(y[m, ] - x)
  }
  return(sum(sum_ll))
}
```

(c) Tempered transition π_0, \dots, π_k , given

$$\pi_0 \sim N\left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, 7 * I_2\right)$$

```
tempLogLikeliHood <- function(x, y, k, K=10)
{
  temp <- (k / K) * logLikeliHood(x, y) + dmvnorm(x, mean=c(0, 0), sigma = 7 * diag(2), log = TRUE)
  return(temp)
}
```

(d) π_n -invariant Metropolis-Hastings (MH) kernel based on a random walk proposal.

```
MHKernel <- function(x, y, k, K=10, theta = 0.5)
{
  # A new proposal
  xProposed = rmvnorm(1, mean = x, sigma = theta^2 * diag(2))
}
```



```

# Acceptance probability
#print(tempLogLikeliHood(xProposed, y, k, K=K))
acc_prob <- tempLogLikeliHood(xProposed, y, k, K=K) -
            tempLogLikeliHood(x, y, k, K=K)
acc_prob <- min(1, exp(acc_prob))
# Sample from uniform to be compared to the acceptance probability
u = runif(1)

# Set next state depending on acceptance probability
if(u <= acc_prob){
  return(xProposed)
}
else{
  return(x)
}
}

```

```
MHKernel(x0, y, 1)
```

```
##           [,1]      [,2]
## [1,] 6.354334 -5.036059
```

```
MHKernel(c(0,0), y, 1)
```

```
##           [,1]      [,2]
## [1,] -0.3577252 -0.2695782
```

(e) Put everything together in an SMC sampler

```

smcSampler <- function(y, K=10, N=100, theta=.5)
{
  # Initialize variables
  particles <- array(0, c(N, K+1, 2))
  ancestorIndices <- matrix(0, nrow = N, ncol = K+1)
  weights <- matrix(1, nrow = N, ncol = K+1)
  normalisedWeights <- matrix(0, nrow = N, ncol = K+1)

  # Initial state
  xInit = rmvnorm(N, mean=c(0,0), 7*diag(2))
  ancestorIndices[, 1] <- 1:N
  particles[, 1, ] <- xInit
  normalisedWeights[, 1] = 1 / N

  for(t in 2:(K+1))
  {
    # Update weights
    for (i in 1:N)
    {
      weights[i, t] = log(normalisedWeights[i, t-1]) +
        tempLogLikeliHood(particles[i, t-1, ], y, k=t, K=K) -
        tempLogLikeliHood(particles[i, t-1, ], y, k=t-1, K=K)
    }
    # normalize weights
    weights[, t] = weights[, t] - max(weights[, t])
    normalisedWeights[, t] = exp(weights[, t]) / sum(exp(weights[, t]))

    newAncestors <- 1:N
    # Resample
    if(1/sum(normalisedWeights[, t-1]^2) < N/2){
      newAncestors <- sample(N, replace = TRUE, prob = normalisedWeights[, t-1])
      normalisedWeights[, t-1] <- 1/N
    }
  }
}

```

```

}

ancestorIndices[, t-1] <- newAncestors

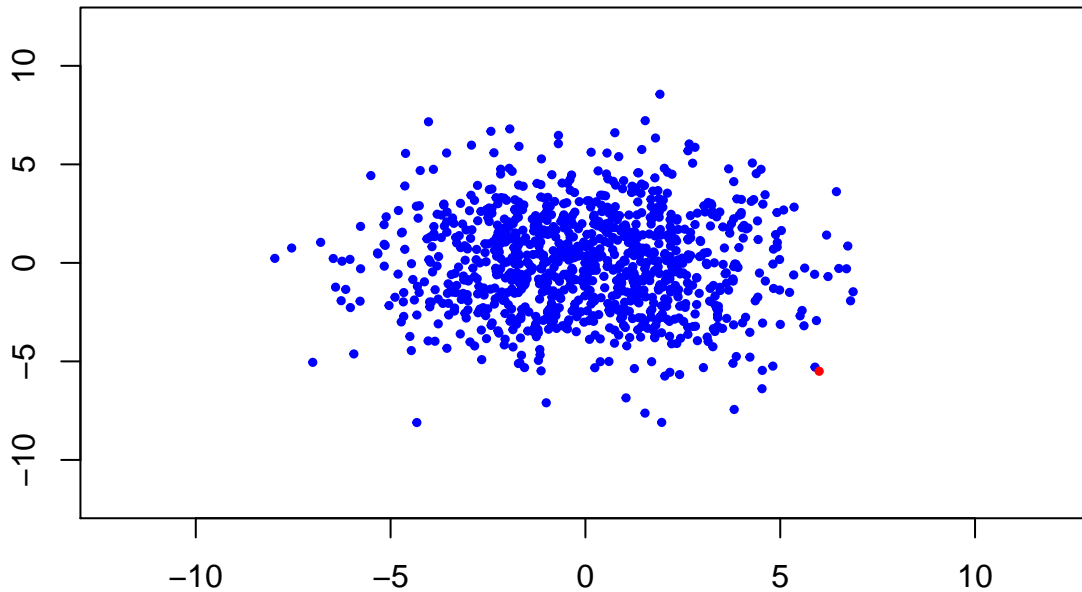
particles[, t, ] <- particles[newAncestors, t-1, ]
# Propagate using MHKernel
for (i in 1:N)
{
  particles[i, t, ] = MHKernel(particles[i, t, ], y,
                               k=t, K=K, theta=theta)
}
}
return(list(particles = particles, weights = normalisedWeights))
}

K=10
result = smcSampler(y, N=1000, K = K)

par(mfrow=c(1, 1))
for (i in 1:11) {
  plot(result$particles[, i, 1], result$particles[, i, 2], xlim = c(-12,12), ylim = c(-12,12), cex = .5, col =
  lines(x0[1], x0[2], col = "red", type = "p", cex = .5, pch = 19)
  print(i)
}

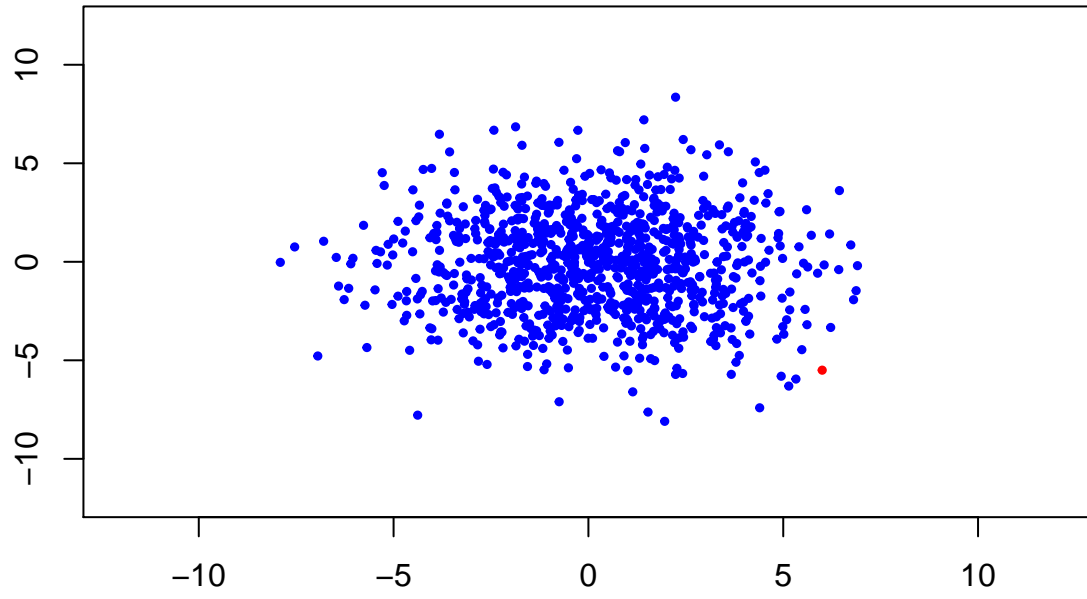
```

1



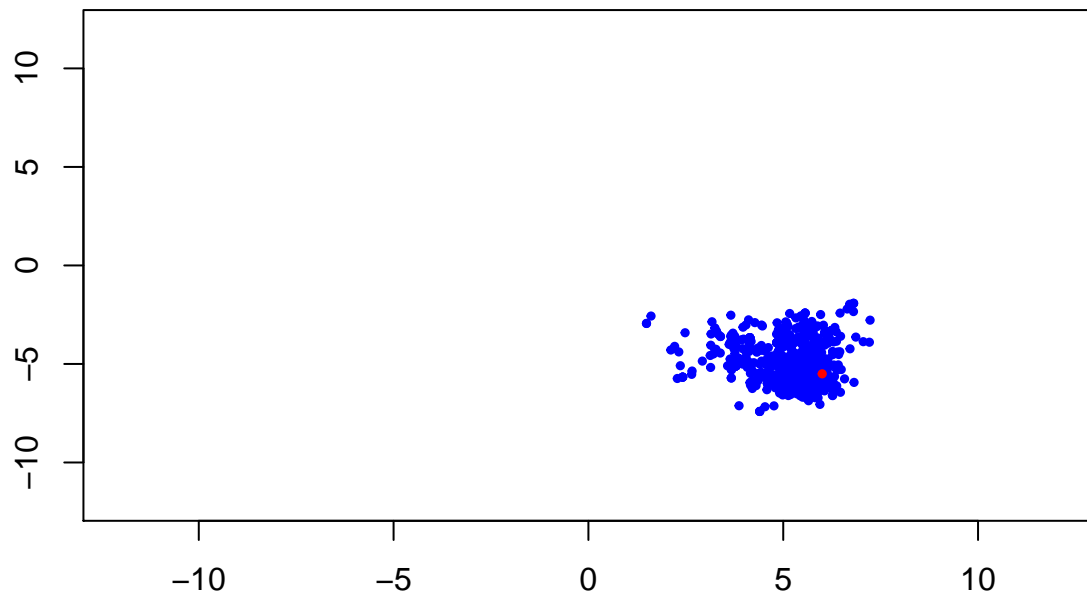
```
## [1] 1
```

2



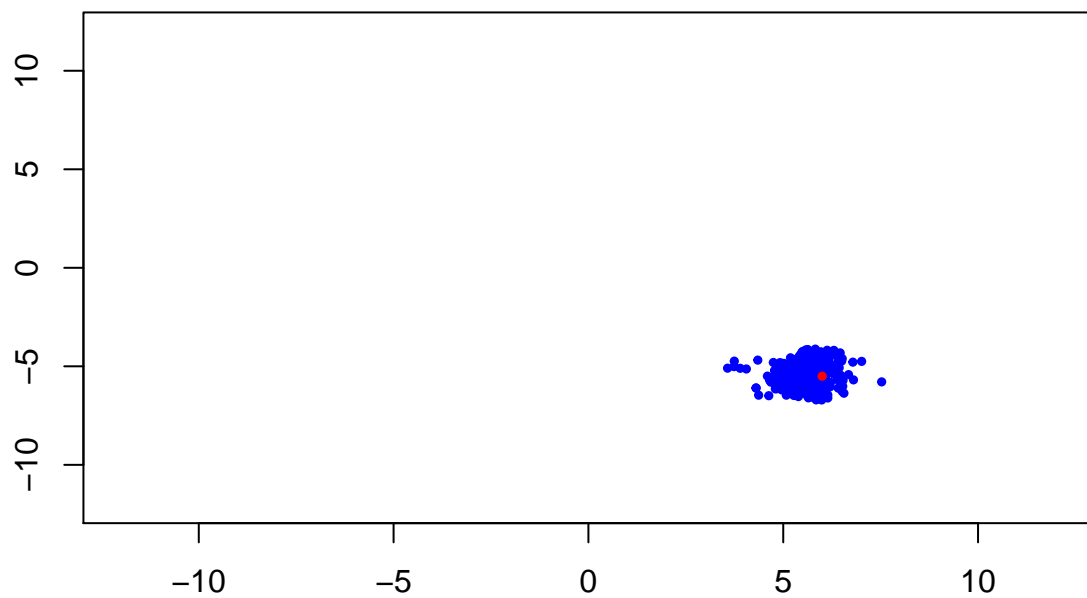
[1] 2

3



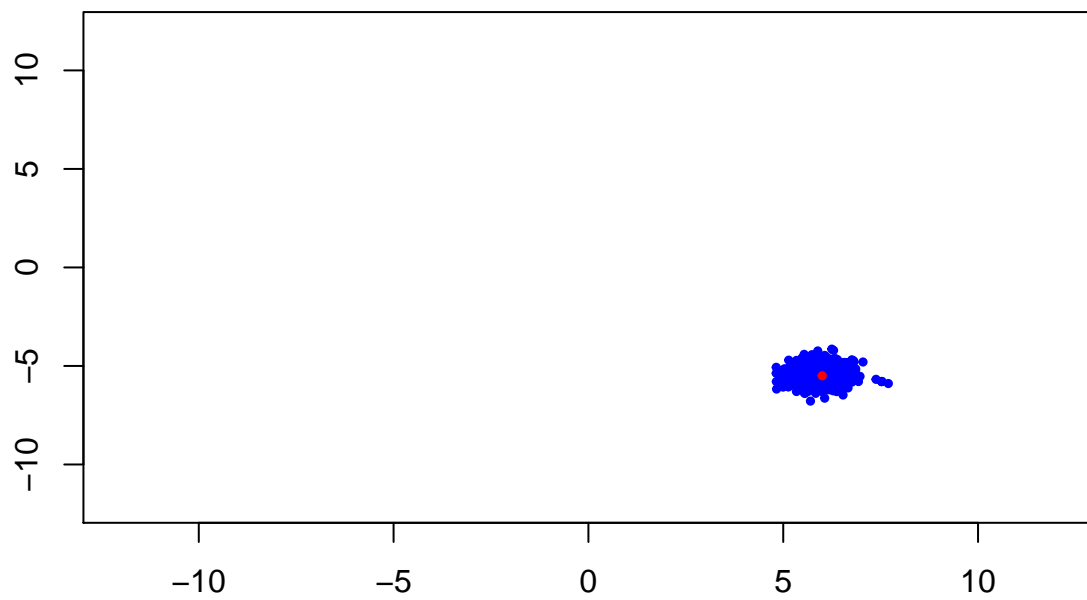
[1] 3

4



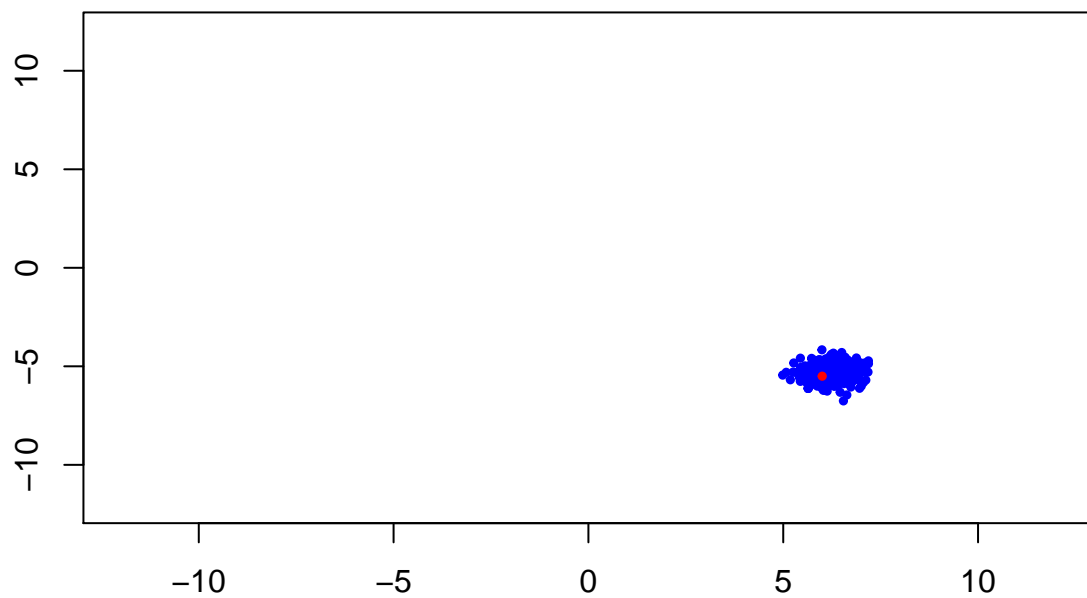
[1] 4

5



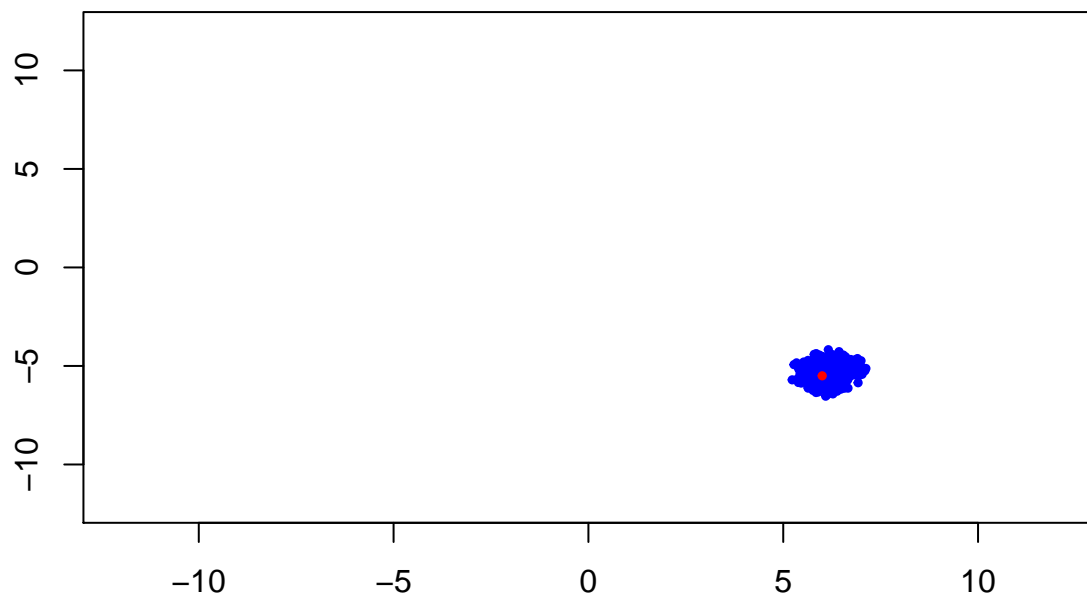
[1] 5

6



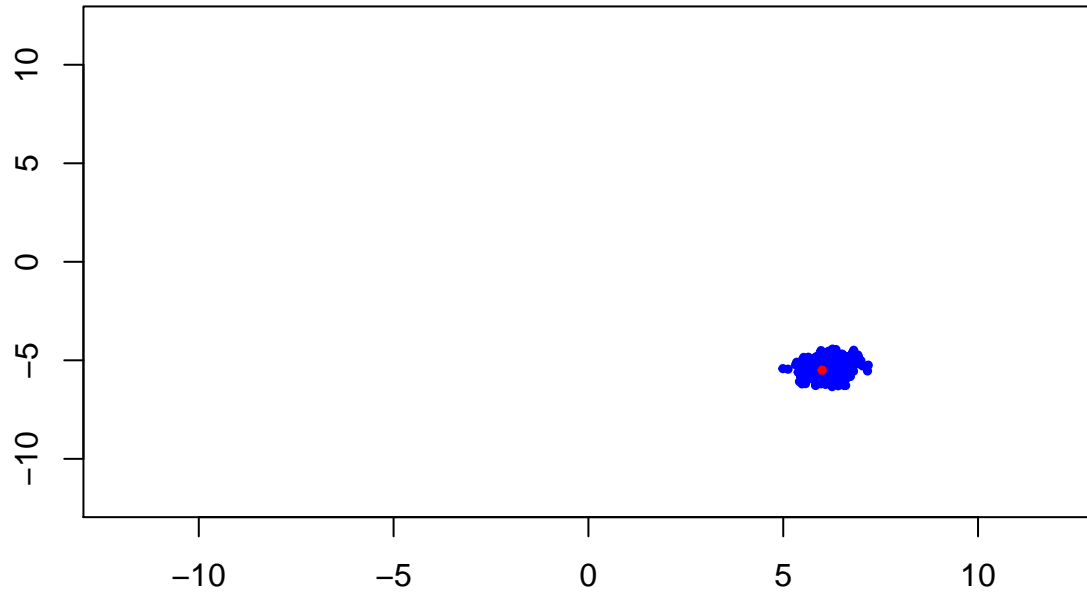
[1] 6

7



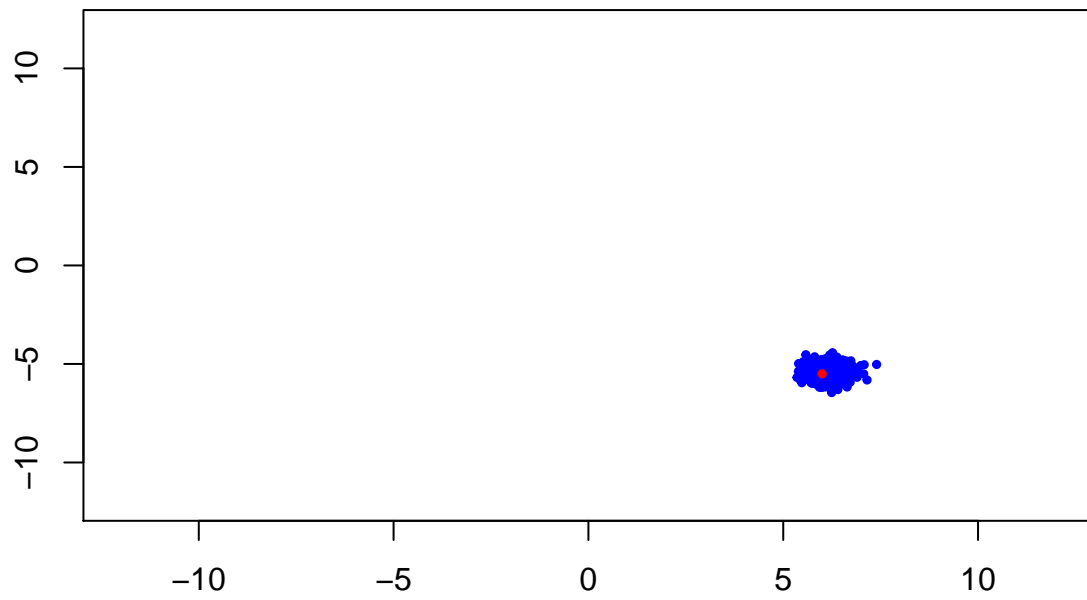
[1] 7

8



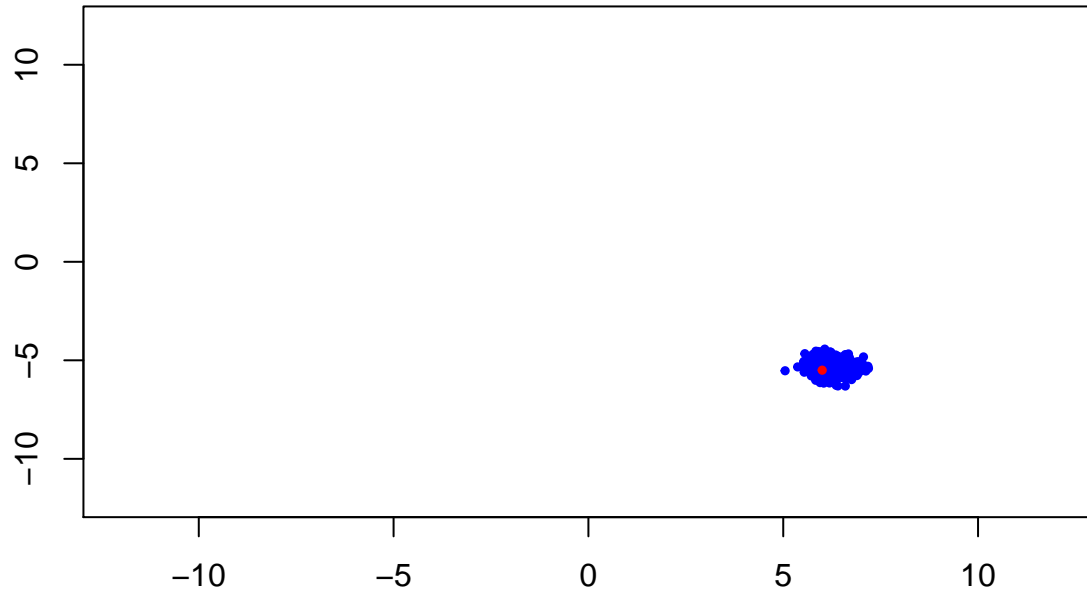
[1] 8

9



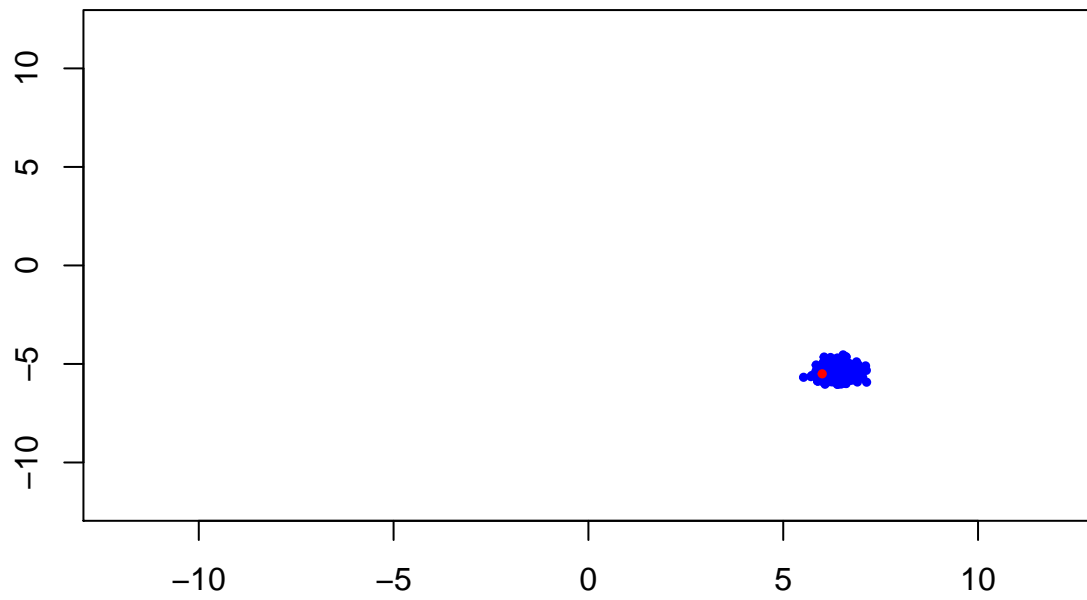
[1] 9

10



[1] 10

11



[1] 11